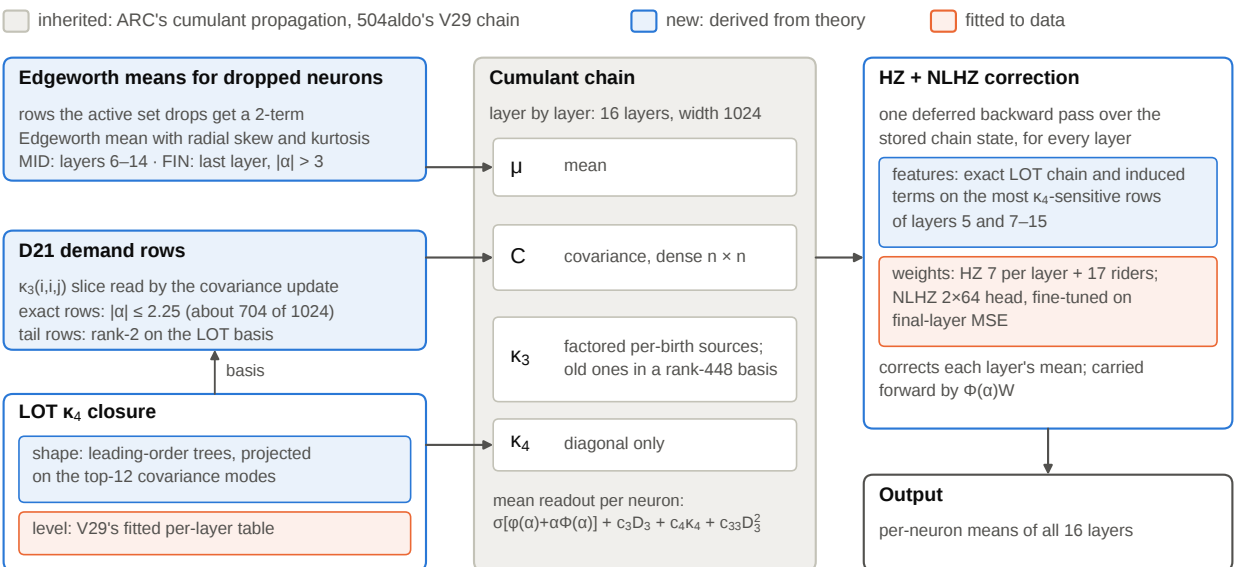
inherited: ARC's cumulant propagation, 504aldo's V29 chain
new: derived from theory
fitted to data
Edgeworth means for dropped neurons
rows the active set drops get a 2-term
Edgeworth mean with radial skew and kurtosis
MID: layers 6–14 · FIN: last layer, |α| > 3
D21 demand rows
κ3(i,i,j) slice read by the covariance update
exact rows: |α| ≤ 2.25 (about 704 of 1024)
tail rows: rank-2 on the LOT basis
basis
LOT κ4 closure
shape: leading-order trees, projected
on the top-12 covariance modes
level: V29's fitted per-layer table
Cumulant chain
layer by layer: 16 layers, width 1024
μ
mean
C
covariance, dense n × n
κ3
factored per-birth sources;
old ones in a rank-448 basis
κ4
diagonal only
mean readout per neuron:
σ[φ(α)+αΦ(α)] + c3D3 + c4κ4 + c33D3²
HZ + NLHZ correction
one deferred backward pass over the
stored chain state, for every layer
features: exact LOT chain and induced
terms on the most κ4-sensitive rows
of layers 5 and 7–15
weights: HZ 7 per layer + 17 riders;
NLHZ 2×64 head, fine-tuned on
final-layer MSE
corrects each layer's mean; carried
forward by Φ(α)W
Output
per-neuron means of all 16 layers